

AI-Native Candidate Discovery & Ranking System

Recruiting at Scale for the Senior AI Engineer (Founding Team) Role

Prepared by: Team Antigravity

Challenge: Redrob Data & AI Challenge 2026

Methodology: Offline Calibrated Heuristics + Multi-stage Filtering

Execution Time: <10 seconds on standard CPU

1. Problem Statement & Constraints

The Recruiting Dilemma: Early-stage Series A startup seeking a Founding AI Engineer with deep ML/systems capability mixed with a shipper mindset.

The Objective: Programmatically score 100,000 resumes and output the **Top 100 candidates**.

Compute Parameter	Hackathon Limit	Antigravity Performance
Runtime Limit	≤ 5 minutes (300 seconds)	~ 3.5 seconds (Fast, linear parsing)
Memory limit	≤ 16 GB RAM	~ 60 MB (Streamed file parsing, zero-memory leaks)
Network Calls	Off (No OpenAI, Gemini, or hosted APIs)	Fully Offline (Self-contained logic)
Hardware	CPU Only — no GPU	Single-threaded CPU

2. Dataset Analysis & Trap Detection

Synthetic recruitment datasets contain systemic noise and intentional 'traps' to penalize naive systems. Our pipeline implements two main filters:

- **Honeypots ($\leq 10\%$ tolerance in top 100):** The dataset contains **84 honeypot profiles** which are structured with impossible data anomalies. If ranked in the top 100, the candidate's actual relevance is 0.
- *Startup Founding Violations:* 76 candidates worked at Krutrim, Sarvam AI, Genpact AI, or Rephrase.ai years before the companies were founded (e.g. at Krutrim in 2019, founded in 2023).
- *Skill Durations:* 8 candidates claimed 'expert' level proficiency in 5+ skills with exactly 0 months of experience.
- **Consulting/Services-Only Trap:** JD explicitly disqualifies consulting/services backgrounds (TCS, Infosys, Wipro, Accenture, Cognizant, etc.). Candidates who have spent their entire careers there are hard-filtered.
- **Role/Title Traps:** Candidates who stuff resumes with AI keywords but hold non-technical current titles (e.g., 'Marketing Manager', 'Accountant') are filtered out.

3. Multi-Stage Pipeline Architecture

Our system streams, filters, scores, and ranks candidates in a clean modular architecture:

Stage	Process Description	Effect / Output
1. Streaming Ingestion	Iterate over candidates.jsonl line-by-line using Python standard json. Zero memory overhead.	Scales to millions of rows seamlessly.
2. Strict Filtration	Verify startup timelines, skill duration anomalies, consulting-only career histories, and current title technical keywords.	Safely excludes all 84 honeypots and non-tech traps.
3. Calibrated Scoring	Compute mathematical scores based on YoE, matching skills (weighted and log-scaled), location bonuses, and behavioral flags.	Produces a continuous score representing JD fit.
4. Sorting & Formatting	Sort candidates by score (descending) and candidate_id (ascending) to break ties. Generate custom reasoning narrations.	Extracts Top 100, formatted to CSV validator specifications.

4. Feature Engineering & Scoring Rules

Valid candidates are scored out of 115 points based on the following breakdown:

Dimension	Rules & Formulas	Max Pts
Experience (YoE)	Ideal: 6–8 years (+15 pts); Target: 5–9 years (+12 pts); Adjacent: 4-5 or 9-11 (+8 pts).	15
Skills Match	Skill score = Skill weight * Proficiency multiplier * Log duration multiplier. - Heavy weight (3.0): RAG, Dense Retrieval, LLMs, Vector DBs (Pinecone, Weaviate). - Medium weight (2.0): Python, Fine-tuning (LoRA), PyTorch, NLP, eval metrics. - Scale: expert=1.0, beginner=0.2. Duration: log-scaled (capped at 3 years).	45
Location Alignment	Noida/Pune (+10 pts); Welcome cities (Delhi/NCR/Mumbai/Hyd) (+5 pts); India + relocate (+5 pts); Outside India (-20 pts).	10
Notice Period	≤ 30 days (+10 pts); ≤ 60 days (+5 pts); 90 days (0 pts); > 90 days (-10 pts).	10
Platform Engagement	Active < 30 days (+5 pts); Inactive > 180 days (-15 pts); Response rate ≥ 80% (+10 pts); Open to work (+5 pts); Github ≥ 60 (+5 pts); Interview Attendance ≥ 80% (+3 pts) or < 50% (-10 pts).	20
Career & Titles	Product company ratio ≥ 75% (+10 pts); Job Hopper average tenure < 18 mo (-8 pts); Current AI/ML Title match (+10 pts) or Software/Backend/Data (+5 pts).	15

5. Results & Top Candidates Profile

Our system processed the complete pool of 100,000 candidates and outputted the top 100 successfully. Here are the leading profiles:

Rank	CID	Name	Current Role	YoE	Location	Score
1	CAND_0061257	Advaith Pillai	Staff ML Engineer @ LinkedIn	8.0	Noida (Preferred)	122.0
2	CAND_0018499	Aarav Trivedi	Senior ML Engineer @ Zomato	7.2	Noida (Preferred)	120.0
3	CAND_0052328	Vikram Banerjee	Recommendation Systems Eng @ Amazon	6.5	Pune (Preferred)	120.0
4	CAND_0007009	Anika Pillai	Recommendation Systems Eng @ Wysa	7.9	Noida (Preferred)	118.0
5	CAND_0026942	Sneha Nair	Junior ML Engineer @ Verloop.io	6.0	Chandigarh (India)	117.0

Key Validation Milestones:

- **Honeypot Rate:** Exactly **0%** in top 100. All 84 honeypot candidates were successfully trapped and filtered.
- **Reasoning Quality:** Custom-generated non-templated text referencing actual candidate details (experience, specific skills, locations, and notice period constraints). No hallucinations.
- **Deterministic Tie-breaking:** Fully validated for lexicographical `candidate\_id` order ascending on score ties.

6. Production Scale & Future Enhancements

To transition this candidate ranking system from offline challenge-heuristics to a production scale talent system, we propose:

1. **Hybrid Retrieval Layer:** Deploy an Elasticsearch index combining BM25 keyword matching (for explicit skill strings) with dense vector embeddings (trained on HR/Resume corpuses like BGE-M3) to capture conceptual fit.
2. **Two-Pass Ranking Architecture (Bi-Encoder + Cross-Encoder):** Use a fast Bi-Encoder to filter 100,000 candidates down to 500 in under a second, then apply a precise Cross-Encoder model (re-ranker) on CPU to score the top 500, maximizing retrieval NDCG without violating latency limits.
3. **Learning-to-Rank (LTR):** Train an XGBoost model (using `LambdaMART` or similar) using historical recruiter clicks, messages, and offer acceptance data to automatically learn feature weights rather than relying on manual coefficients.
4. **Dynamic Embedding Refresh:** Build an asynchronous pipeline using Kafka and Pinecone/Qdrant to instantly re-index profiles when candidates update their skills, preventing index drift.
5. **Automated Evaluation Harness:** Set up offline validation using NDCG and MAP over historical test sets, and implement split-traffic A/B testing in production to monitor real-world click-through rates.